REPORT

Merged results, rounds 1 and 2

Historical generated tables (21 October 2026).

Download PDF

Back to the decision brief

Contract: prompts `a579532922ca`, scorers `d438bf722d3a`, cases v1 `d92bcbc12d72`, confirm `a8e2b349e391`. Model labels are `model@side [variant]`. Failed cases count as zero. No model graded another model.

Chart questions (ChartQA, relaxed accuracy)

MODEL	CASES	FAILED	CORRECT	CORRECT_HUMAN	CORRECT_AUGMENTED	MEDIAN S
ollama:qwen2.5vl:32b@desktop	200	0	82.0%	75.0%	89.0%	0.880
ollama:qwen2.5vl:32b@desktop [expert]	200	0	82.0%	79.0%	85.0%	0.900
ollama:qwen2.5vl:32b@desktop [selfcheck]	200	0	87.0%	84.0%	90.0%	1.610
ollama:qwen2.5vl:32b@desktop [upscale]	200	1	76.5%	73.0%	80.0%	1.580
ollama:qwen2.5vl:7b@desktop	200	0	53.5%	63.0%	44.0%	0.460
ollama:qwen2.5vl:7b@desktop [expert]	200	0	44.5%	53.0%	36.0%	0.480
ollama:qwen2.5vl:7b@desktop [selfcheck]	200	1	80.0%	75.0%	85.0%	0.640
ollama:qwen2.5vl:7b@desktop [upscale]	200	0	49.0%	63.0%	35.0%	0.860
ollama:qwen2.5vl:7b@laptop	200	0	52.5%	62.0%	43.0%	0.880
ollama:qwen2.5vl:7b@laptop [expert]	200	0	47.5%	56.0%	39.0%	0.900
ollama:qwen2.5vl:7b@laptop [selfcheck]	200	1	79.0%	74.0%	84.0%	1.480
ollama:qwen2.5vl:7b@laptop [upscale]	200	0	47.5%	62.0%	33.0%	2.190
ollama:qwen3-vl:8b-instruct@desktop	200	0	78.0%	64.0%	92.0%	0.520
ollama:qwen3-vl:8b-instruct@desktop [expert]	200	0	76.0%	74.0%	78.0%	0.530
ollama:qwen3-vl:8b-instruct@desktop [selfcheck]	200	0	32.5%	46.0%	19.0%	0.710
ollama:qwen3-vl:8b-instruct@desktop [upscale]	200	0	79.5%	65.0%	94.0%	0.790
ollama:qwen3-vl:8b-instruct@laptop	200	0	79.5%	67.0%	92.0%	1.430
ollama:qwen3-vl:8b-instruct@laptop [expert]	200	0	75.0%	73.0%	77.0%	1.090
ollama:qwen3-vl:8b-instruct@laptop [selfcheck]	200	0	36.0%	50.0%	22.0%	1.650
ollama:qwen3-vl:8b-instruct@laptop [upscale]	200	0	80.0%	66.0%	94.0%	2.130
ollama:qwen3.8:27b@desktop	200	0	85.5%	77.0%	94.0%	0.640
ollama:qwen3.8:27b@desktop [expert]	200	0	87.0%	81.0%	93.0%	0.680

MODEL	CASES	FAILED	CORRECT	CORRECT_HUMAN	CORRECT_AUGMENTED	MEDIAN S
ollama:qwen3.8:27b@desktop [hires]	200	0	85.5%	77.0%	94.0%	0.670
ollama:qwen3.8:27b@desktop [selfcheck]	200	0	39.5%	41.0%	38.0%	0.960
ollama:qwen3.8:27b@desktop [upscale]	200	0	88.0%	81.0%	95.0%	1.330
openai:qwen3vl32b@desktop	200	0	85.0%	75.0%	95.0%	0.520
openai:qwen3vl32b@desktop [expert]	200	0	84.0%	75.0%	93.0%	0.480
openai:qwen3vl32b@desktop [hires]	200	0	85.0%	75.0%	95.0%	0.500
openai:qwen3vl32b@desktop [selfcheck]	200	0	91.0%	88.0%	94.0%	1.400
openai:qwen3vl32b@desktop [table only]	200	0	81.0%	70.0%	92.0%	0.360
openai:qwen3vl32b@desktop [upscale]	200	0	89.5%	84.0%	95.0%	1.090

Chart questions, held-out confirmation set

MODEL	CASES	FAILED	CORRECT	CORRECT_HUMAN	CORRECT_AUGMENTED	MEDIAN S
ollama:qwen3.8:27b@desktop	2022	0	83.3%	71.7%	94.5%	0.640
openai:qwen3vl32b@desktop	2022	0	86.4%	77.9%	94.5%	0.490
openai:qwen3vl32b@desktop [tuned]	2022	0	88.3%	83.3%	93.0%	1.460

Table transcription (TEDS, cell F1)

MODEL	CASES	FAILED	TEDS	TEDS_STRUCTURE	CELL_F1	NUMERIC_RECALL	MEDIAN S
ollama:qwen2.5vl:32b@desktop	60	0	0.805	0.849	0.789	0.972	5.070
ollama:qwen2.5vl:32b@desktop [expert]	60	0	0.807	0.869	0.801	0.967	4.950
ollama:qwen2.5vl:32b@desktop [selfcheck]	60	0	0.770	0.817	0.790	0.967	5.830
ollama:qwen2.5vl:32b@desktop [upscale]	60	0	0.830	0.872	0.797	0.973	5.230
ollama:qwen2.5vl:7b@desktop	60	0	0.724	0.790	0.761	0.933	1.520
ollama:qwen2.5vl:7b@desktop [expert]	60	0	0.748	0.812	0.766	0.951	1.570
ollama:qwen2.5vl:7b@desktop [selfcheck]	60	0	0.752	0.820	0.767	0.940	1.830
ollama:qwen2.5vl:7b@desktop [upscale]	60	0	0.734	0.799	0.778	0.938	1.820
ollama:qwen2.5vl:7b@laptop	60	0	0.724	0.797	0.765	0.941	5.250
ollama:qwen2.5vl:7b@laptop [expert]	60	0	0.757	0.817	0.767	0.953	5.440
ollama:qwen2.5vl:7b@laptop [selfcheck]	60	0	0.749	0.817	0.765	0.939	5.780
ollama:qwen2.5vl:7b@laptop [upscale]	60	0	0.728	0.792	0.775	0.925	6.750
ollama:qwen3-vl:8b-instruct@desktop	60	0	0.819	0.862	0.788	0.950	1.880
ollama:qwen3-vl:8b-instruct@desktop [expert]	60	0	0.822	0.869	0.789	0.948	1.890
ollama:qwen3-vl:8b-instruct@desktop [selfcheck]	60	0	0.818	0.866	0.790	0.945	2.000
ollama:qwen3-vl:8b-instruct@desktop [upscale]	60	0	0.832	0.873	0.794	0.947	1.990

MODEL	CASES	FAILED	TEDS	TEDS_STRUCTURE	CELL_F1	NUMERIC_RECALL	MEDIAN S
ollama:qwen3-vl:8b-instruct@laptop	60	0	0.818	0.861	0.788	0.944	8.540
ollama:qwen3-vl:8b-instruct@laptop [expert]	60	0	0.829	0.870	0.789	0.946	6.550
ollama:qwen3-vl:8b-instruct@laptop [selfcheck]	60	0	0.823	0.873	0.790	0.946	6.320
ollama:qwen3-vl:8b-instruct@laptop [upscale]	60	0	0.831	0.873	0.793	0.946	7.010
ollama:qwen3.8:27b@desktop	60	0	0.852	0.896	0.777	0.968	2.640
ollama:qwen3.8:27b@desktop [expert]	60	0	0.870	0.912	0.796	0.975	2.500
ollama:qwen3.8:27b@desktop [hires]	60	0	0.853	0.896	0.777	0.969	2.650
ollama:qwen3.8:27b@desktop [selfcheck]	60	0	0.847	0.892	0.792	0.971	2.580
ollama:qwen3.8:27b@desktop [upscale]	60	0	0.852	0.891	0.784	0.977	3.370
openai:qwen3vl32b@desktop	60	0	0.821	0.861	0.776	0.969	5.510
openai:qwen3vl32b@desktop [expert]	60	0	0.849	0.891	0.782	0.973	5.330
openai:qwen3vl32b@desktop [hires]	60	0	0.821	0.861	0.776	0.969	5.520
openai:qwen3vl32b@desktop [selfcheck]	60	0	0.828	0.872	0.779	0.955	5.700
openai:qwen3vl32b@desktop [upscale]	60	0	0.841	0.885	0.787	0.970	5.560

Table transcription, held-out confirmation set

MODEL	CASES	FAILED	TEDS	TEDS_STRUCTURE	CELL_F1	NUMERIC_RECALL	MEDIAN S
ollama:qwen3.8:27b@desktop	159	1	0.870	0.913	0.772	0.948	2.650
openai:qwen3vl32b@desktop	159	0	0.851	0.907	0.766	0.967	5.640
openai:qwen3vl32b@desktop [tuned]	159	0	0.841	0.893	0.759	0.967	5.650

Chart data extraction (values within 5%)

MODEL	CASES	FAILED	VALUE_RECALL	VALUE_PRECISION	LABEL_RECALL	MEDIAN S
ollama:qwen2.5vl:32b@desktop	60	0	0.721	0.720	0.949	3.290
ollama:qwen2.5vl:32b@desktop [expert]	60	0	0.723	0.735	0.846	3.730
ollama:qwen2.5vl:32b@desktop [selfcheck]	60	0	0.742	0.729	0.935	4.200
ollama:qwen2.5vl:32b@desktop [upscale]	60	0	0.683	0.677	0.944	4.000
ollama:qwen2.5vl:7b@desktop	60	0	0.715	0.700	0.607	0.990
ollama:qwen2.5vl:7b@desktop [expert]	60	3	0.533	0.516	0.578	1.040
ollama:qwen2.5vl:7b@desktop [selfcheck]	60	0	0.708	0.671	0.829	1.110
ollama:qwen2.5vl:7b@desktop [upscale]	60	3	0.784	0.783	0.593	1.430
ollama:qwen2.5vl:7b@laptop	60	2	0.650	0.635	0.587	2.580
ollama:qwen2.5vl:7b@laptop [expert]	60	4	0.541	0.535	0.579	2.860
ollama:qwen2.5vl:7b@laptop [selfcheck]	60	0	0.626	0.620	0.850	2.790

MODEL	CASES	FAILED	VALUE_RECALL	VALUE_PRECISION	LABEL_RECALL	MEDIAN S
ollama:qwen2.5vl:7b@laptop [upscale]	60	2	0.636	0.627	0.676	3.870
ollama:qwen3-vl:8b-instruct@desktop	60	0	0.483	0.494	0.417	1.000
ollama:qwen3-vl:8b-instruct@desktop [expert]	60	0	0.915	0.921	0.804	1.500
ollama:qwen3-vl:8b-instruct@desktop [selfcheck]	60	0	0.930	0.959	0.921	1.610
ollama:qwen3-vl:8b-instruct@desktop [upscale]	60	0	0.643	0.653	0.593	1.520
ollama:qwen3-vl:8b-instruct@laptop	60	0	0.418	0.427	0.350	1.750
ollama:qwen3-vl:8b-instruct@laptop [expert]	60	0	0.900	0.899	0.816	4.180
ollama:qwen3-vl:8b-instruct@laptop [selfcheck]	60	0	0.930	0.958	0.921	4.420
ollama:qwen3-vl:8b-instruct@laptop [upscale]	60	0	0.630	0.637	0.583	4.490
ollama:qwen3.8:27b@desktop	60	0	0.864	0.859	0.963	2.100
ollama:qwen3.8:27b@desktop [expert]	60	1	0.730	0.721	0.963	2.050
ollama:qwen3.8:27b@desktop [hires]	60	0	0.864	0.859	0.963	2.160
ollama:qwen3.8:27b@desktop [selfcheck]	60	0	0.817	0.825	0.979	2.240
ollama:qwen3.8:27b@desktop [upscale]	60	2	0.884	0.881	0.944	2.720
openai:qwen3vl32b@desktop	60	0	0.901	0.906	0.912	3.940
openai:qwen3vl32b@desktop [expert]	60	0	0.893	0.906	0.929	3.750
openai:qwen3vl32b@desktop [hires]	60	0	0.901	0.906	0.912	3.780
openai:qwen3vl32b@desktop [selfcheck]	60	0	0.902	0.909	0.947	4.110
openai:qwen3vl32b@desktop [upscale]	60	0	0.902	0.916	0.905	4.600

Prose crops (character error rate, lower is better)

MODEL	CASES	FAILED	CER_COMPLETED	NUMERIC_RECALL	NUMERIC_PRECISION	MEDIAN S
ollama:qwen2.5vl:32b@desktop	50	0	0.012	0.965	0.953	1.450
ollama:qwen2.5vl:7b@desktop	50	0	0.018	0.965	0.953	0.510
ollama:qwen2.5vl:7b@laptop	50	0	0.020	0.965	0.953	1.760
ollama:qwen3-vl:8b-instruct@desktop	50	0	0.012	0.965	0.953	0.610
ollama:qwen3-vl:8b-instruct@laptop	50	0	0.012	0.965	0.953	2.560
ollama:qwen3.8:27b@desktop	50	0	0.011	0.965	0.953	0.710
openai:qwen3vl32b@desktop	50	0	0.012	0.965	0.953	1.000

Whole pages (word F1)

MODEL	CASES	FAILED	WORD_F1	NUMERIC_RECALL	MEDIAN S
openai:qwen3vl32b@desktop	40	0	0.809	0.909	16.500
openai:qwen3vl32b@desktop [layoutmask]	40	0	0.895	0.720	11.870

Layout (figures and tables, IoU >= 0.5)

MODEL	FIGURE P	FIGURE R	FIGURE F1	TABLE P	TABLE R	TABLE F1
openai:qwen3vl32b@desktop	0.570	0.658	0.611	0.600	0.882	0.714
yolo@desktop	0.836	0.872	0.854	0.782	0.843	0.811

Variants against baseline (paired, same cases)

MODEL	TASK	VARIANT	CASES	BASELINE	VARIANT	CHANGE	95% CI	BETTER/WORSE	TEST	P	FAMILY
ollama:qwen2.5vl:32b@desktop	chart_qa	expert	200	0.82	0.82	+0.0000	[-0.04, 0.04]	9/9	exact McNemar	1.0	exploratory
ollama:qwen2.5vl:32b@desktop	chart_qa	selfcheck	200	0.82	0.87	+0.0500	[0.0, 0.1]	18/8	exact McNemar	0.0755	exploratory
ollama:qwen2.5vl:32b@desktop	chart_qa	upscale	200	0.82	0.765	-0.0550	[-0.105, 0.0]	9/20	exact McNemar	0.0614	exploratory
ollama:qwen2.5vl:32b@desktop	chart_table	expert	60	0.7214	0.7234	+0.0020	[-0.1149, 0.1169]	10/11	paired bootstrap	0.966	exploratory
ollama:qwen2.5vl:32b@desktop	chart_table	selfcheck	60	0.7214	0.7417	+0.0203	[-0.0781, 0.1174]	9/6	paired bootstrap	0.66	exploratory
ollama:qwen2.5vl:32b@desktop	chart_table	upscale	60	0.7214	0.6826	-0.0388	[-0.1011, 0.0127]	4/5	paired bootstrap	0.157	exploratory
ollama:qwen2.5vl:32b@desktop	table	expert	60	0.8046	0.8073	+0.0027	[-0.0326, 0.0331]	14/11	paired bootstrap	0.804	exploratory
ollama:qwen2.5vl:32b@desktop	table	selfcheck	60	0.8046	0.7697	-0.0350	[-0.0781, 0.0083]	9/27	paired bootstrap	0.111	exploratory
ollama:qwen2.5vl:32b@desktop	table	upscale	60	0.8046	0.8304	+0.0258	[-0.0007, 0.0585]	17/8	paired bootstrap	0.062	exploratory
ollama:qwen2.5vl:7b@desktop	chart_qa	expert	200	0.535	0.445	-0.0900	[-0.135, -0.04]	4/22	exact McNemar	0.0005	exploratory
ollama:qwen2.5vl:7b@desktop	chart_qa	selfcheck	200	0.535	0.8	+0.2650	[0.19, 0.34]	63/10	exact McNemar	0.0	exploratory
ollama:qwen2.5vl:7b@desktop	chart_qa	upscale	200	0.535	0.49	-0.0450	[-0.085, -0.005]	5/14	exact McNemar	0.0636	exploratory
ollama:qwen2.5vl:7b@desktop	chart_table	expert	60	0.7153	0.5328	-0.1825	[-0.318, -0.0478]	5/19	paired bootstrap	0.005	exploratory
ollama:qwen2.5vl:7b@desktop	chart_table	selfcheck	60	0.7153	0.7078	-0.0075	[-0.1192, 0.1044]	8/9	paired bootstrap	0.907	exploratory
ollama:qwen2.5vl:7b@desktop	chart_table	upscale	60	0.7153	0.7837	+0.0684	[0.011, 0.1389]	7/4	paired bootstrap	0.017	exploratory
ollama:qwen2.5vl:7b@desktop	table	expert	60	0.724	0.7476	+0.0236	[-0.0114, 0.0649]	21/11	paired bootstrap	0.223	exploratory
ollama:qwen2.5vl:7b@desktop	table	selfcheck	60	0.724	0.7524	+0.0285	[-0.0206, 0.0813]	18/20	paired bootstrap	0.252	exploratory

MODEL	TASK	VARIANT	CASES	BASELINE	VARIANT	CHANGE	95% CI	BETTER/WORSE	TEST	P	FAMILY
ollama:qwen2.5vl:7b@desktop	table	upscale	60	0.724	0.7338	+0.0098	[-0.0212, 0.0478]	11/14	paired bootstrap	0.635	exploratory
ollama:qwen2.5vl:7b@laptop	chart_qa	expert	200	0.525	0.475	-0.0500	[-0.09, -0.01]	4/14	exact McNemar	0.0309	exploratory
ollama:qwen2.5vl:7b@laptop	chart_qa	selfcheck	200	0.525	0.79	+0.2650	[0.185, 0.34]	64/11	exact McNemar	0.0	exploratory
ollama:qwen2.5vl:7b@laptop	chart_qa	upscale	200	0.525	0.475	-0.0500	[-0.095, -0.01]	5/15	exact McNemar	0.0414	exploratory
ollama:qwen2.5vl:7b@laptop	chart_qa_confirm2	selfcheck	1322	0.5575	0.7352	+0.1778	[0.1498, 0.2073]	324/89	exact McNemar	0.0	supporting_v
ollama:qwen2.5vl:7b@laptop	chart_qa_confirm2	upscale	1322	0.5575	0.5197	-0.0378	[-0.0552, -0.0212]	44/94	exact McNemar	0.0	supporting_v
ollama:qwen2.5vl:7b@laptop	chart_qa_confirm2	vote3	1322	0.5575	0.6029	+0.0454	[0.0303, 0.0605]	80/20	exact McNemar	0.0	confirmatory_
ollama:qwen2.5vl:7b@laptop	chart_qa_dev2	selfcheck	586	0.558	0.7696	+0.2116	[0.1672, 0.2543]	162/38	exact McNemar	0.0	exploratory_v
ollama:qwen2.5vl:7b@laptop	chart_qa_dev2	upscale	586	0.558	0.5102	-0.0478	[-0.0734, -0.0222]	14/42	exact McNemar	0.0002	exploratory_v
ollama:qwen2.5vl:7b@laptop	chart_qa_dev2	vote3	586	0.558	0.6007	+0.0427	[0.0205, 0.0648]	36/11	exact McNemar	0.0003	exploratory_v
ollama:qwen2.5vl:7b@laptop	chart_table	expert	60	0.6497	0.5406	-0.1091	[-0.2478, 0.0364]	7/20	paired bootstrap	0.116	exploratory
ollama:qwen2.5vl:7b@laptop	chart_table	selfcheck	60	0.6497	0.6256	-0.0241	[-0.1643, 0.115]	11/15	paired bootstrap	0.678	exploratory
ollama:qwen2.5vl:7b@laptop	chart_table	upscale	60	0.6497	0.6359	-0.0138	[-0.1296, 0.1017]	10/10	paired bootstrap	0.826	exploratory
ollama:qwen2.5vl:7b@laptop	page_text_confirm2	layoutmask	80	0.7574	0.8428	+0.0854	[0.0385, 0.135]	39/31	paired bootstrap	0.0	confirmatory_
ollama:qwen2.5vl:7b@laptop	table	expert	60	0.7239	0.7574	+0.0335	[-0.0001, 0.0704]	22/10	paired bootstrap	0.052	exploratory
ollama:qwen2.5vl:7b@laptop	table	selfcheck	60	0.7239	0.7491	+0.0252	[-0.0247, 0.0779]	19/20	paired bootstrap	0.317	exploratory
ollama:qwen2.5vl:7b@laptop	table	upscale	60	0.7239	0.7284	+0.0045	[-0.0225, 0.0332]	12/12	paired bootstrap	0.757	exploratory
ollama:qwen3-vl:8b-instruct@desktop	chart_qa	expert	200	0.78	0.76	-0.0200	[-0.075, 0.035]	13/17	exact McNemar	0.5847	exploratory
ollama:qwen3-vl:8b-instruct@desktop	chart_qa	selfcheck	200	0.78	0.325	-0.4550	[-0.535, -0.375]	9/100	exact McNemar	0.0	exploratory
ollama:qwen3-vl:8b-instruct@desktop	chart_qa	upscale	200	0.78	0.795	+0.0150	[-0.01, 0.04]	5/2	exact McNemar	0.4531	exploratory
ollama:qwen3-vl:8b-instruct@desktop	chart_table	expert	60	0.4827	0.9154	+0.4327	[0.3014, 0.5664]	28/3	paired bootstrap	0.0	exploratory

MODEL	TASK	VARIANT	CASES	BASELINE	VARIANT	CHANGE	95% CI	BETTER/WORSE	TEST	P	FAMILY
ollama:qwen3-vl:8b-instruct@desktop	chart_table	selfcheck	60	0.4827	0.9302	+0.4475	[0.3265, 0.5708]	29/3	paired bootstrap	0.0	exploratory
ollama:qwen3-vl:8b-instruct@desktop	chart_table	upscale	60	0.4827	0.6429	+0.1602	[0.0581, 0.2755]	11/5	paired bootstrap	0.005	exploratory
ollama:qwen3-vl:8b-instruct@desktop	table	expert	60	0.8186	0.8218	+0.0032	[-0.0178, 0.022]	9/6	paired bootstrap	0.757	exploratory
ollama:qwen3-vl:8b-instruct@desktop	table	selfcheck	60	0.8186	0.8181	-0.0005	[-0.0155, 0.0158]	15/9	paired bootstrap	0.946	exploratory
ollama:qwen3-vl:8b-instruct@desktop	table	upscale	60	0.8186	0.8321	+0.0135	[-0.0008, 0.0399]	6/4	paired bootstrap	0.142	exploratory
ollama:qwen3-vl:8b-instruct@laptop	chart_qa	expert	200	0.795	0.75	-0.0450	[-0.095, 0.005]	9/18	exact McNemar	0.1221	exploratory
ollama:qwen3-vl:8b-instruct@laptop	chart_qa	selfcheck	200	0.795	0.36	-0.4350	[-0.51, -0.35]	9/96	exact McNemar	0.0	exploratory
ollama:qwen3-vl:8b-instruct@laptop	chart_qa	upscale	200	0.795	0.8	+0.0050	[-0.025, 0.035]	5/4	exact McNemar	1.0	exploratory
ollama:qwen3-vl:8b-instruct@laptop	chart_qa_confirm2	selfcheck	1322	0.7496	0.323	-0.4266	[-0.4554, -0.3964]	50/614	exact McNemar	0.0	supporting_v
ollama:qwen3-vl:8b-instruct@laptop	chart_qa_confirm2	upscale	1322	0.7496	0.7648	+0.0151	[0.003, 0.0272]	44/24	exact McNemar	0.0205	supporting_v
ollama:qwen3-vl:8b-instruct@laptop	chart_qa_confirm2	vote3	1322	0.7496	0.7481	-0.0015	[-0.0106, 0.0076]	19/21	exact McNemar	0.8746	confirmatory_
ollama:qwen3-vl:8b-instruct@laptop	chart_qa_dev2	selfcheck	586	0.785	0.3225	-0.4625	[-0.5068, -0.4164]	18/289	exact McNemar	0.0	exploratory_v
ollama:qwen3-vl:8b-instruct@laptop	chart_qa_dev2	upscale	586	0.785	0.7935	+0.0085	[-0.0102, 0.0273]	18/13	exact McNemar	0.4731	exploratory_v
ollama:qwen3-vl:8b-instruct@laptop	chart_qa_dev2	vote3	586	0.785	0.7799	-0.0051	[-0.0171, 0.0068]	5/8	exact McNemar	0.5811	exploratory_v
ollama:qwen3-vl:8b-instruct@laptop	chart_table	expert	60	0.4179	0.8999	+0.4820	[0.349, 0.5997]	29/1	paired bootstrap	0.0	exploratory
ollama:qwen3-vl:8b-instruct@laptop	chart_table	selfcheck	60	0.4179	0.9296	+0.5118	[0.3876, 0.6394]	33/2	paired bootstrap	0.0	exploratory
ollama:qwen3-vl:8b-instruct@laptop	chart_table	upscale	60	0.4179	0.6304	+0.2125	[0.0982, 0.3287]	15/3	paired bootstrap	0.001	exploratory
ollama:qwen3-vl:8b-instruct@laptop	page_text_confirm2	layoutmask	80	0.8283	0.838	+0.0098	[-0.0354, 0.0504]	34/35	paired bootstrap	0.66	confirmatory_
ollama:qwen3-vl:8b-instruct@laptop	table	expert	60	0.8177	0.8285	+0.0109	[0.0007, 0.0254]	9/4	paired bootstrap	0.018	exploratory
ollama:qwen3-vl:8b-instruct@laptop	table	selfcheck	60	0.8177	0.8231	+0.0054	[-0.0174, 0.0299]	13/8	paired bootstrap	0.686	exploratory
ollama:qwen3-vl:8b-instruct@laptop	table	upscale	60	0.8177	0.8312	+0.0136	[-0.0006, 0.0392]	8/4	paired bootstrap	0.128	exploratory

MODEL	TASK	VARIANT	CASES	BASELINE	VARIANT	CHANGE	95% CI	BETTER/WORSE	TEST	P	FAMILY
ollama:qwen3.8:27b@desktop	chart_qa	expert	200	0.855	0.87	+0.0150	[-0.02, 0.055]	9/6	exact McNemar	0.6072	exploratory
ollama:qwen3.8:27b@desktop	chart_qa	hires	200	0.855	0.855	+0.0000	[0.0, 0.0]	0/0	exact McNemar	1.0	exploratory
ollama:qwen3.8:27b@desktop	chart_qa	selfcheck	200	0.855	0.395	-0.4600	[-0.53, -0.385]	3/95	exact McNemar	0.0	exploratory
ollama:qwen3.8:27b@desktop	chart_qa	upscale	200	0.855	0.88	+0.0250	[-0.015, 0.07]	12/7	exact McNemar	0.3593	exploratory
ollama:qwen3.8:27b@desktop	chart_table	expert	60	0.8639	0.7299	-0.1340	[-0.2507, -0.0179]	4/12	paired bootstrap	0.018	exploratory
ollama:qwen3.8:27b@desktop	chart_table	hires	60	0.8639	0.8639	+0.0000	[0.0, 0.0]	0/0	paired bootstrap	1.0	exploratory
ollama:qwen3.8:27b@desktop	chart_table	selfcheck	60	0.8639	0.817	-0.0469	[-0.1296, 0.0239]	3/7	paired bootstrap	0.183	exploratory
ollama:qwen3.8:27b@desktop	chart_table	upscale	60	0.8639	0.8837	+0.0198	[-0.0495, 0.1003]	6/2	paired bootstrap	0.526	exploratory
ollama:qwen3.8:27b@desktop	table	expert	60	0.8525	0.8695	+0.0170	[-0.0088, 0.0518]	11/9	paired bootstrap	0.296	exploratory
ollama:qwen3.8:27b@desktop	table	hires	60	0.8525	0.8529	+0.0004	[-0.0, 0.0012]	2/1	paired bootstrap	0.285	exploratory
ollama:qwen3.8:27b@desktop	table	selfcheck	60	0.8525	0.8474	-0.0051	[-0.0342, 0.0262]	11/15	paired bootstrap	0.676	exploratory
ollama:qwen3.8:27b@desktop	table	upscale	60	0.8525	0.8524	-0.0001	[-0.01, 0.0095]	18/9	paired bootstrap	0.944	exploratory
openai:qwen3vl32b@desktop	chart_qa	expert	200	0.85	0.84	-0.0100	[-0.04, 0.015]	3/5	exact McNemar	0.7266	exploratory
openai:qwen3vl32b@desktop	chart_qa	hires	200	0.85	0.85	+0.0000	[0.0, 0.0]	0/0	exact McNemar	1.0	exploratory
openai:qwen3vl32b@desktop	chart_qa	selfcheck	200	0.85	0.91	+0.0600	[0.015, 0.105]	17/5	exact McNemar	0.0169	exploratory
openai:qwen3vl32b@desktop	chart_qa	upscale	200	0.85	0.895	+0.0450	[0.01, 0.08]	11/2	exact McNemar	0.0225	exploratory
openai:qwen3vl32b@desktop	chart_qa_confirm	tuned	2022	0.864	0.8828	+0.0188	[0.0059, 0.0321]	113/75	exact McNemar	0.0068	confirmatory
openai:qwen3vl32b@desktop	chart_qa_confirm2	selfcheck	1322	0.82	0.8449	+0.0250	[0.0091, 0.0416]	78/45	exact McNemar	0.0037	supporting_v
openai:qwen3vl32b@desktop	chart_qa_confirm2	upscale	1322	0.82	0.8351	+0.0151	[0.0023, 0.0272]	45/25	exact McNemar	0.0225	supporting_v
openai:qwen3vl32b@desktop	chart_qa_confirm2	vote3	1322	0.82	0.8457	+0.0257	[0.0136, 0.0371]	49/15	exact McNemar	0.0	confirmatory_
openai:qwen3vl32b@desktop	chart_qa_dev2	readfirst	586	0.8379	0.8567	+0.0188	[-0.0051, 0.0444]	34/23	exact McNemar	0.1849	exploratory_v

MODEL	TASK	VARIANT	CASES	BASELINE	VARIANT	CHANGE	95% CI	BETTER/WORSE	TEST	P	FAMILY
openai:qwen3vl32b@desktop	chart_qa_dev2	selfcheck	586	0.8379	0.8652	+0.0273	[0.0051, 0.0495]	30/14	exact McNemar	0.0226	exploratory_v
openai:qwen3vl32b@desktop	chart_qa_dev2	upscale	586	0.8379	0.8584	+0.0205	[0.0017, 0.0392]	21/9	exact McNemar	0.0428	exploratory_v
openai:qwen3vl32b@desktop	chart_qa_dev2	vote3	586	0.8379	0.8669	+0.0290	[0.0137, 0.0461]	19/2	exact McNemar	0.0002	exploratory_v
openai:qwen3vl32b@desktop	chart_table	expert	60	0.9009	0.8931	-0.0078	[-0.0486, 0.0336]	1/4	paired bootstrap	0.669	exploratory
openai:qwen3vl32b@desktop	chart_table	hires	60	0.9009	0.9009	+0.0000	[0.0, 0.0]	0/0	paired bootstrap	1.0	exploratory
openai:qwen3vl32b@desktop	chart_table	selfcheck	60	0.9009	0.9017	+0.0008	[-0.0492, 0.0508]	2/2	paired bootstrap	0.899	exploratory
openai:qwen3vl32b@desktop	chart_table	upscale	60	0.9009	0.9021	+0.0012	[-0.0487, 0.0499]	3/2	paired bootstrap	0.888	exploratory
openai:qwen3vl32b@desktop	page_text	layoutmask	40	0.8094	0.8951	+0.0858	[0.0403, 0.1353]	20/16	paired bootstrap	0.0	exploratory_v
openai:qwen3vl32b@desktop	page_text_confirm2	layoutmask	80	0.8138	0.8805	+0.0668	[0.033, 0.101]	37/36	paired bootstrap	0.0	confirmatory_
openai:qwen3vl32b@desktop	table	expert	60	0.8212	0.8486	+0.0274	[0.0017, 0.0594]	10/11	paired bootstrap	0.034	exploratory
openai:qwen3vl32b@desktop	table	hires	60	0.8212	0.8213	+0.0001	[-0.0, 0.0004]	2/1	paired bootstrap	0.51	exploratory
openai:qwen3vl32b@desktop	table	selfcheck	60	0.8212	0.8278	+0.0067	[-0.0197, 0.0353]	15/15	paired bootstrap	0.654	exploratory
openai:qwen3vl32b@desktop	table	upscale	60	0.8212	0.8409	+0.0197	[-0.0019, 0.0444]	15/12	paired bootstrap	0.075	exploratory
openai:qwen3vl32b@desktop	table_confirm	tuned	159	0.8508	0.8414	-0.0094	[-0.0295, 0.0099]	24/33	paired bootstrap	0.323	confirmatory

Reproducibility: the same model run on both machines (laptop minus desktop)

MODEL	TASK	VARIANT	SHARED CASES	SAME WEIGHTS	IDENTICAL ANSWERS	IDENTICAL SCORES	DESKTOP	LAPTOP	DIFFERENCE [95% CI]	VERDICT
ollama:qwen2.5vl:7b	chart_qa	baseline	200	yes	94.5%	98.0%	0.535	0.525	-0.0100 [-0.03, 0.01]	reproduced
ollama:qwen2.5vl:7b	chart_qa	expert	200	yes	91.5%	96.0%	0.445	0.475	+0.0300 [0.005, 0.06]	sides disagree
ollama:qwen2.5vl:7b	chart_qa	selfcheck	200	yes	64.5%	94.0%	0.8	0.79	-0.0100 [-0.045, 0.025]	reproduced
ollama:qwen2.5vl:7b	chart_qa	upscale	200	yes	93.0%	97.5%	0.49	0.475	-0.0150 [-0.04, 0.005]	reproduced
ollama:qwen2.5vl:7b	chart_table	baseline	60	yes	66.7%	86.7%	0.7153	0.6497	-0.0656 [-0.1444, 0.0067]	reproduced

MODEL	TASK	VARIANT	SHARED CASES	SAME WEIGHTS	IDENTICAL ANSWERS	IDENTICAL SCORES	DESKTOP	LAPTOP	DIFFERENCE [95% CI]	VERDICT
ollama:qwen2.5vl:7b	chart_table	expert	60	yes	60.0%	93.3%	0.5328	0.5406	+0.0079 [-0.0463, 0.0625]	reproduced
ollama:qwen2.5vl:7b	chart_table	selfcheck	60	yes	38.3%	73.3%	0.7078	0.6256	-0.0821 [-0.1886, 0.0193]	reproduced
ollama:qwen2.5vl:7b	chart_table	upscale	60	yes	55.0%	75.0%	0.7837	0.6359	-0.1477 [-0.2549, -0.0593]	sides disagree
ollama:qwen2.5vl:7b	ocr_crop	baseline	50	yes	90.0%	96.0%	0.0243	0.0253	+0.0011 [0.0, 0.0034]	reproduced
ollama:qwen2.5vl:7b	table	baseline	60	yes	80.0%	83.3%	0.724	0.7239	-0.0000 [-0.0103, 0.0108]	reproduced
ollama:qwen2.5vl:7b	table	expert	60	yes	71.7%	75.0%	0.7476	0.7574	+0.0098 [-0.0018, 0.0239]	reproduced
ollama:qwen2.5vl:7b	table	selfcheck	60	yes	76.7%	83.3%	0.7524	0.7491	-0.0033 [-0.0091, 0.0014]	reproduced
ollama:qwen2.5vl:7b	table	upscale	60	yes	88.3%	88.3%	0.7338	0.7284	-0.0054 [-0.0193, 0.0028]	reproduced
ollama:qwen3-vl:8b-instruct	chart_qa	baseline	200	NO	93.0%	97.5%	0.78	0.795	+0.0150 [-0.005, 0.035]	different weights: not a reproduction test
ollama:qwen3-vl:8b-instruct	chart_qa	expert	200	NO	92.0%	96.0%	0.76	0.75	-0.0100 [-0.035, 0.015]	different weights: not a reproduction test
ollama:qwen3-vl:8b-instruct	chart_qa	selfcheck	200	NO	51.0%	95.5%	0.325	0.36	+0.0350 [0.01, 0.065]	different weights: not a reproduction test
ollama:qwen3-vl:8b-instruct	chart_qa	upscale	200	NO	96.0%	98.5%	0.795	0.8	+0.0050 [-0.01, 0.02]	different weights: not a reproduction test
ollama:qwen3-vl:8b-instruct	chart_table	baseline	60	NO	70.0%	88.3%	0.4827	0.4179	-0.0648 [-0.1463, 0.0037]	different weights: not a reproduction test
ollama:qwen3-vl:8b-instruct	chart_table	expert	60	NO	75.0%	98.3%	0.9154	0.8999	-0.0155 [-0.0466, 0.0]	different weights: not a reproduction test
ollama:qwen3-vl:8b-instruct	chart_table	selfcheck	60	NO	83.3%	98.3%	0.9302	0.9296	-0.0006 [-0.0017, 0.0]	different weights: not a reproduction test
ollama:qwen3-vl:8b-instruct	chart_table	upscale	60	NO	65.0%	83.3%	0.6429	0.6304	-0.0125 [-0.1125, 0.0833]	different weights: not a reproduction test
ollama:qwen3-vl:8b-instruct	ocr_crop	baseline	50	NO	96.0%	96.0%	0.0123	0.0124	+0.0001 [0.0, 0.0003]	different weights: not a reproduction test
ollama:qwen3-vl:8b-instruct	table	baseline	60	NO	80.0%	83.3%	0.8186	0.8177	-0.0009 [-0.0025, 0.0004]	different weights: not a reproduction test
ollama:qwen3-vl:8b-instruct	table	expert	60	NO	80.0%	86.7%	0.8218	0.8285	+0.0068 [-0.0017, 0.0221]	different weights: not a reproduction test
ollama:qwen3-vl:8b-instruct	table	selfcheck	60	NO	78.3%	91.7%	0.8181	0.8231	+0.0049 [-0.012, 0.0245]	different weights: not a reproduction test
ollama:qwen3-vl:8b-instruct	table	upscale	60	NO	88.3%	91.7%	0.8321	0.8312	-0.0009 [-0.0026, 0.0]	different weights: not a reproduction test